



Here's a detailed explanation for all your questions with examples and step-by-step explanations. Let's cover each question in detail, starting with table creation, optimization techniques, and other related concepts.

1. Database Architecture for Express Cargo Module

Entities: Flight, Cargo, Pricing, and Revenue

To design a database for an express cargo module, the entities and their relationships must be well-defined.

Tables and Relationships:

- 1. **Flight Table:**
 - Stores flight details.
 - `FlightID` is the primary key.
- 2. **Cargo Table:**
 - Stores cargo details, linked to flights.
 - `CargoID` is the primary key.
 - `FlightID` is a foreign key referencing `FlightID` in the `Flight` table.
- 3. **Pricing Table:**
 - Stores pricing details for cargo types.
 - `PricingID` is the primary key.
- 4. **Revenue Table:**
 - Tracks revenue for flights based on cargo bookings.
 - `RevenueID` is the primary key.

Schema for Tables

1. Flight Table:

```
sql

CREATE TABLE Flight (
  FlightID INT AUTO INCREMENT PRIMARY KEY,
  FlightNumber VARCHAR(20) NOT NULL,
  Destination VARCHAR(50),
  DepartureTime DATETIME
);
```

2. Cargo Table:

```
sql

CREATE TABLE Cargo (
  CargoID INT AUTO INCREMENT PRIMARY KEY,
  CargoType VARCHAR(50),
  Weight DECIMAL(10, 2),
```

```
Volume DECIMAL(10, 2),
FlightID INT,
FOREIGN KEY (FlightID) REFERENCES Flight(FlightID)
);
```

3. Pricing Table:

sql

```
CREATE TABLE Pricing (
  PricingID INT AUTO_INCREMENT PRIMARY KEY,
  CargoType VARCHAR(50),
  PricePerKg DECIMAL(10, 2),
  PricePerVolume DECIMAL(10, 2)
);
```

4. Revenue Table:

sql

```
CREATE TABLE Revenue (
  RevenueID INT AUTO_INCREMENT PRIMARY KEY,
  FlightID INT,
  TotalRevenue DECIMAL(10, 2),
  FOREIGN KEY (FlightID) REFERENCES Flight(FlightID)
);
```

Table Relationships:

- **Flight → Cargo (One-to-Many):** A flight can carry multiple cargos.
- **Pricing → Cargo (One-to-Many):** Each cargo type has a specific pricing.
- **Flight → Revenue (One-to-One):** Each flight has one revenue record.

2. Techniques to Optimize Database Operations

1. Indexing:

- Speeds up query performance.
- Example: Index the `FlightNumber` column for quick searches.

sql

```
CREATE INDEX idx_flight_number ON Flight(FlightNumber);
```

2. Partitioning:

- Divides a table into smaller, manageable pieces.
- Example: Partition `Flight` data by `Destination`.

3. Normalization:

- Reduces data redundancy by splitting tables.

4. Denormalization:

- Combines frequently joined tables to reduce query time.

5. Caching:

- Stores frequently accessed data in memory.

6. Query Optimization:

- Use `EXPLAIN` to analyze queries.

```
sql
```

```
EXPLAIN SELECT * FROM Cargo WHERE Weight > 100;
```

3. Database Partitioning

Partitioning splits a large table into smaller, manageable chunks based on a column, such as `DATE` or `LOCATION`. This helps with:

- **Query Performance:** Only relevant partitions are scanned.
- **Maintenance:** Easier to manage subsets of data.

Example:

Partition the `Flight` table by `Destination`:

```
sql
```

```
CREATE TABLE Flight (
  FlightID INT AUTO_INCREMENT PRIMARY KEY,
  FlightNumber VARCHAR(20) NOT NULL,
  Destination VARCHAR(50),
  DepartureTime DATETIME
) PARTITION BY LIST COLUMNS (Destination) (
  PARTITION p_north VALUES IN ('NYC', 'BOS'),
  PARTITION p_west VALUES IN ('SFO', 'LAX')
);
```

4. Table Schema with Primary and Foreign Keys

1. Primary Key:

- Uniquely identifies a record in a table.
- Example: `FlightID` in the `Flight` table.

2. Foreign Key:

- Links records in one table to another.
- Example: `FlightID` in the `Cargo` table references `FlightID` in the `Flight` table.

Example Schema:

```
sql
```

```
CREATE TABLE Flight (
  FlightID INT AUTO_INCREMENT PRIMARY KEY,
  FlightNumber VARCHAR(20) NOT NULL
);

CREATE TABLE Cargo (
  CargoID INT AUTO_INCREMENT PRIMARY KEY,
  FlightID INT,
  FOREIGN KEY (FlightID) REFERENCES Flight(FlightID)
);
```

5. Index a Query to Find Maximum Flight Number for a Destination

Query:

```
sql
```

```
SELECT MAX(FlightNumber) FROM Flight WHERE Destination = 'NYC';
```

Index for Optimization:

```
sql
```

```
CREATE INDEX idx_destination ON Flight(Destination);
```

6. Select MAX(FlightNumber) for 2 Million Records with Group By

Query:

```
sql
```

```
SELECT Destination, MAX(FlightNumber)
FROM Flight
WHERE DepartureTime > '2024-01-01'
GROUP BY Destination;
```

Optimization Techniques:

1. Index on `DepartureTime`:

```
sql
```

```
CREATE INDEX idx_departure_time ON Flight(DepartureTime);
```

2. Partitioning by `Destination`.

7. Indexing in MySQL

What Are Indexes? Indexes are data structures that improve query performance by allowing faster lookups.

Example:

Index the `FlightNumber` column:

```
sql

CREATE INDEX idx_flight_number ON Flight(FlightNumber);
```

How Indexing Works:

Internally, MySQL uses **B-trees** or **hash tables** for indexes. It creates a sorted data structure, reducing the number of rows scanned during a query.

8. Indexing vs. Triggers

Feature	Indexing	Triggers
Purpose	Speeds up query performance.	Automates actions based on events.
Usage	Improves `SELECT` queries.	Used for enforcing rules or logging changes.
Example	Index on `FlightNumber`.	Trigger on `INSERT` to log changes.

9. Database Transactions

What Are Transactions? A **transaction** is a group of operations performed as a single unit of work, ensuring data consistency.

Properties (ACID):

1. **Atomicity:** All operations succeed or none.
2. **Consistency:** Maintains database integrity.
3. **Isolation:** Transactions don't interfere with each other.
4. **Durability:** Changes persist even after system failures.

Example:

```
sql

START TRANSACTION;
UPDATE Seats SET Available = Available - 1 WHERE FlightID = 101;
```