

Answer:

Normalize the schema and index fields used in filters.

Steps to Design:

1. Use normalized tables for better data organization.
2. Create indexes on columns frequently used in search filters.
3. Use dynamic queries or a search engine like Elasticsearch for complex filters.

Code Example:

sql

```
CREATE INDEX idx_product_category_price ON products(category, price);
SELECT * FROM products
WHERE category = 'Electronics' AND price BETWEEN 100 AND 500;
```

Explanation:

- Indexing reduces filter query times, and search engines can handle unstructured or full-text searches.

**Q33: You



put separator and then add heading



Database Scenario-Based Questions

Q31: A report needs to display data from multiple tables with millions of rows. How would you optimize the query for performance?

Answer:

Use query optimization techniques like indexing, denormalization, and partitioning.

Steps to Optimize:

1. **Indexes:** Add composite or covering indexes to speed up join conditions.
2. **Partitioning:** Partition large tables to reduce query scanning time.
3. **Materialized Views:** Use materialized views for precomputed aggregated data.
4. **Query Caching:** Cache frequently accessed report data.

Code Example:

sql

```
CREATE INDEX idx_report_date ON orders(order_date);
CREATE MATERIALIZED VIEW report_summary AS
SELECT customer_id, SUM(order_amount) AS total_spent
FROM orders
GROUP BY customer_id;
```

Explanation:

- Indexing speeds up table joins and filtering. Materialized views reduce redundant computations.

Q32: An application must allow users to search with complex filters. How would you design the database schema and queries?

Answer:

Normalize the schema and index fields used in filters.

Steps to Design:

1. Use normalized tables for better data organization.
2. Create indexes on columns frequently used in search filters.
3. Use dynamic queries or a search engine like Elasticsearch for complex filters.

Code Example:

sql

```
CREATE INDEX idx_product_category_price ON products(category, price);
SELECT * FROM products
WHERE category = 'Electronics' AND price BETWEEN 100 AND 500;
```

Explanation:

- Indexing reduces filter query times, and search engines can handle unstructured or full-text searches.

Q33: You notice deadlocks in the database during high traffic. How would you detect and fix them?

Answer:

Deadlocks occur when two or more transactions block each other.

Steps to Fix:

1. **Detect:** Use database logs or tools (e.g., `SHOW ENGINE INNODB STATUS` in MySQL).
2. **Resolve:**
 - Break large transactions into smaller ones.
 - Use consistent ordering for accessing tables in queries.
 - Set a timeout for transactions.

Code Example:

```
sql

SET innodb_lock_wait_timeout = 50; -- Set transaction timeout
```

Explanation:

- Analyzing deadlock logs helps identify conflicting transactions, and consistent ordering prevents circular waiting.

Q34: Your database is running out of connections under load. How would you configure and tune connection pooling?

Answer:

Use connection pooling tools like HikariCP in Spring Boot to manage database connections efficiently.

Steps to Configure:

1. **Limit Connections:** Configure `maximumPoolSize` in HikariCP to prevent overloading.
2. **Timeout Settings:** Set connection and idle timeout values.
3. **Monitor Usage:** Use database monitoring tools to optimize pool size.

Code Example:

```
properties

spring.datasource.hikari.maximum-pool-size=20
spring.datasource.hikari.connection-timeout=30000
spring.datasource.hikari.idle-timeout=60000
```

Explanation:

- Connection pooling reuses existing connections instead of creating new ones, saving time and resources.

Q35: You need to migrate a production database to a new schema with minimal downtime. How would you plan this?

Answer:

Plan and execute a zero-downtime migration strategy.

Steps to Migrate:

1. **Versioning:** Use tools like Liquibase or Flyway for schema versioning.
2. **Blue-Green Deployment:** Deploy the new schema in parallel with the old one.
3. **Data Synchronization:** Migrate data incrementally using database triggers or scripts.

Code Example:

```
yaml
```

```
# Flyway migration script
CREATE TABLE new_table (
  id SERIAL PRIMARY KEY,
  name VARCHAR(100)
);
```

Explanation:

- Versioning ensures schema changes are applied consistently, and blue-green deployment minimizes downtime.

Q36: A table has a high number of updates and inserts, leading to degraded performance. How would you handle this?

Answer:

Optimize table structure and database settings to handle high write operations.

Steps to Optimize:

1. **Partitioning:** Split the table into smaller, manageable partitions.
2. **Index Management:** Remove unused indexes that slow down inserts and updates.
3. **Buffering:** Use in-memory buffers like Redis to handle frequent writes.

Code Example:

```
sql
```

```
CREATE TABLE orders PARTITION BY RANGE (order_date) (
  PARTITION p2024 VALUES LESS THAN ('2024-01-01'),
  PARTITION p2025 VALUES LESS THAN ('2025-01-01')
);
```

Explanation:

- Partitioning reduces the scope of operations, improving query performance for high-transaction tables.

Q37: Your application supports soft deletes but query performance is deteriorating. How would you resolve this?

Answer:

Use indexed flags and archival strategies.

Steps to Optimize:

1. **Index Soft Delete Flag:** Add an index to the `is_deleted` column.
2. **Archival:** Move deleted records to a separate archive table periodically.
3. **Filtered Indexes:** Create indexes that exclude soft-deleted records.

Code Example:

```
sql

CREATE INDEX idx_active_records ON users(is_deleted)
WHERE is_deleted = false;
```

Explanation:

- Indexing ensures filtering on `is_deleted` is efficient, and archival reduces table size.

Q38: Multiple microservices need to access the same database but must not interfere with each other's data. How would you handle this?

Answer:

Isolate data for each microservice or use shared data carefully.

Steps to Handle:

1. **Schema Per Service:** Create separate schemas for each microservice.
2. **Access Control:** Use database roles and permissions.
3. **Event Sourcing:** Use an event-driven architecture for shared data access.

Code Example:

```
sql

CREATE SCHEMA service_a AUTHORIZATION user_a;
CREATE SCHEMA service_b AUTHORIZATION user_b;
```

Explanation:

- Isolating schemas avoids interference, and event sourcing ensures consistent updates across services.

Q39: You need to maintain an audit trail of all data changes in a database. How would you implement this?

Answer:

Implement audit logging using triggers or ORM frameworks.

Steps to Implement:

1. **Database Triggers:** Use triggers to log changes to an audit table.
2. **Hibernate Envers:** Use Hibernate Envers for versioning and auditing.
3. **Custom Audit Tables:** Create tables to store audit logs.

Code Example:

```
sql
```

```
CREATE TRIGGER audit_log AFTER UPDATE ON users
FOR EACH ROW
INSERT INTO audit_logs (user_id, old_value, new_value)
VALUES (OLD.id, OLD.name, NEW.name);
```

Explanation:

- Triggers capture changes automatically, and Envers provides ORM-based auditing.

Q40: Your application is moving from a relational database to a NoSQL database. How would you handle this migration?

Answer:

Plan data model transformation and incremental migration.

Steps to Migrate:

1. **Data Model:** Map relational tables to NoSQL collections.
2. **Incremental Migration:** Migrate data in phases to ensure correctness.
3. **Testing:** Validate data consistency post-migration.

Code Example:

```
json
```

```
{
  "user_id": 1,
  "name": "John Doe",
  "orders": [
    {"order_id": 101, "amount": 250},
    {"order_id": 102, "amount": 450}
  ]
}
```

Explanation:

- NoSQL databases use document or key-value structures, simplifying access patterns.

Let me know if you need further adjustments or clarifications!